ROYAL INSTITUTE
OF TECHNOLOGY

Networked Systems Security

Domain Name System Security

Panos Papadimitratos
Networked Systems Security Group
www.ee.kth.se/nss

ROYAL INSTITUTE OF TECHNOLOGY

Domain Name System (DNS)

- Translate names to IP addresses
  - www.ee.kth.se q 130.237.45.45
- End-hosts query DNS servers for name-to-IP translation
- Requests are handled in a recursive manner
  - Tree-like structure

- Note: DNS was not designed with security in mind

Tree structure of domain names

- Internet root
- Top-level domain servers
- Second-level domain names, registered by organizations (or individuals)
- Sub-domain names created by such organizations
- Names assigned to hosts
- E.g., www.ee.kth.se/nss

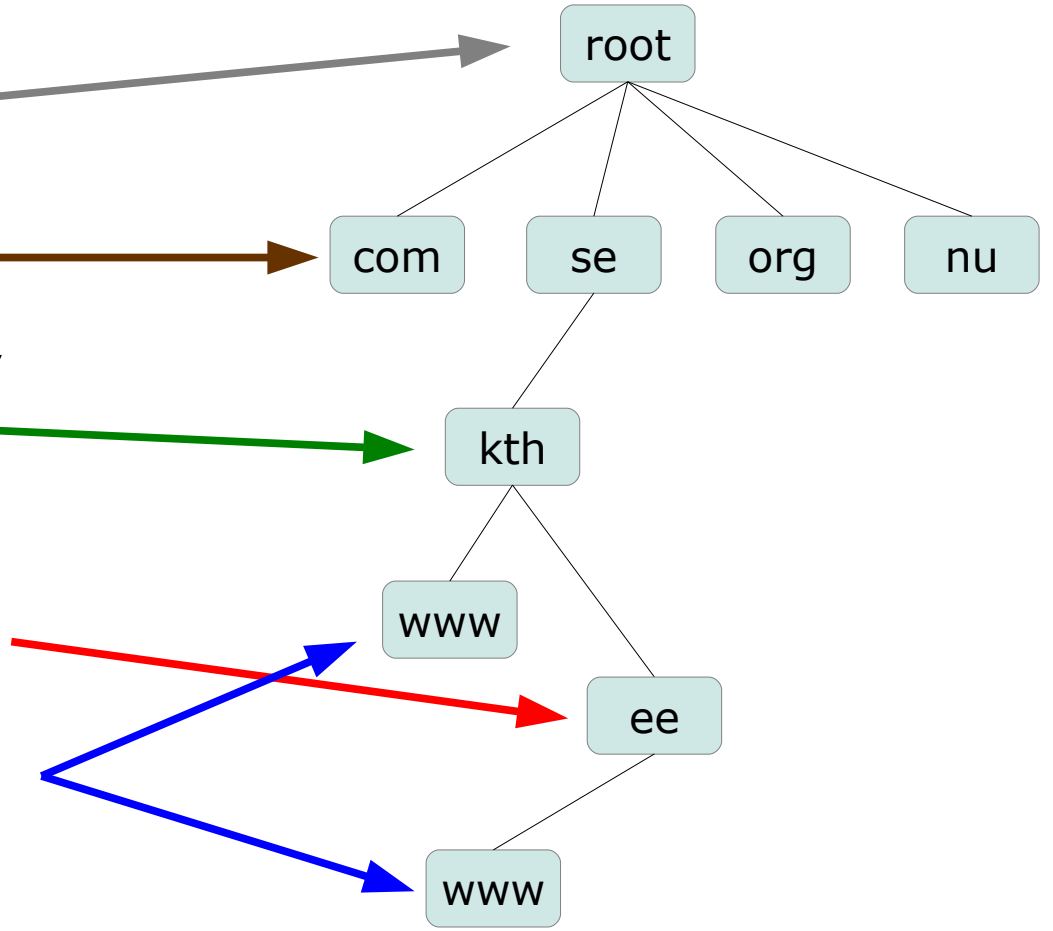

Tree structure of authorities

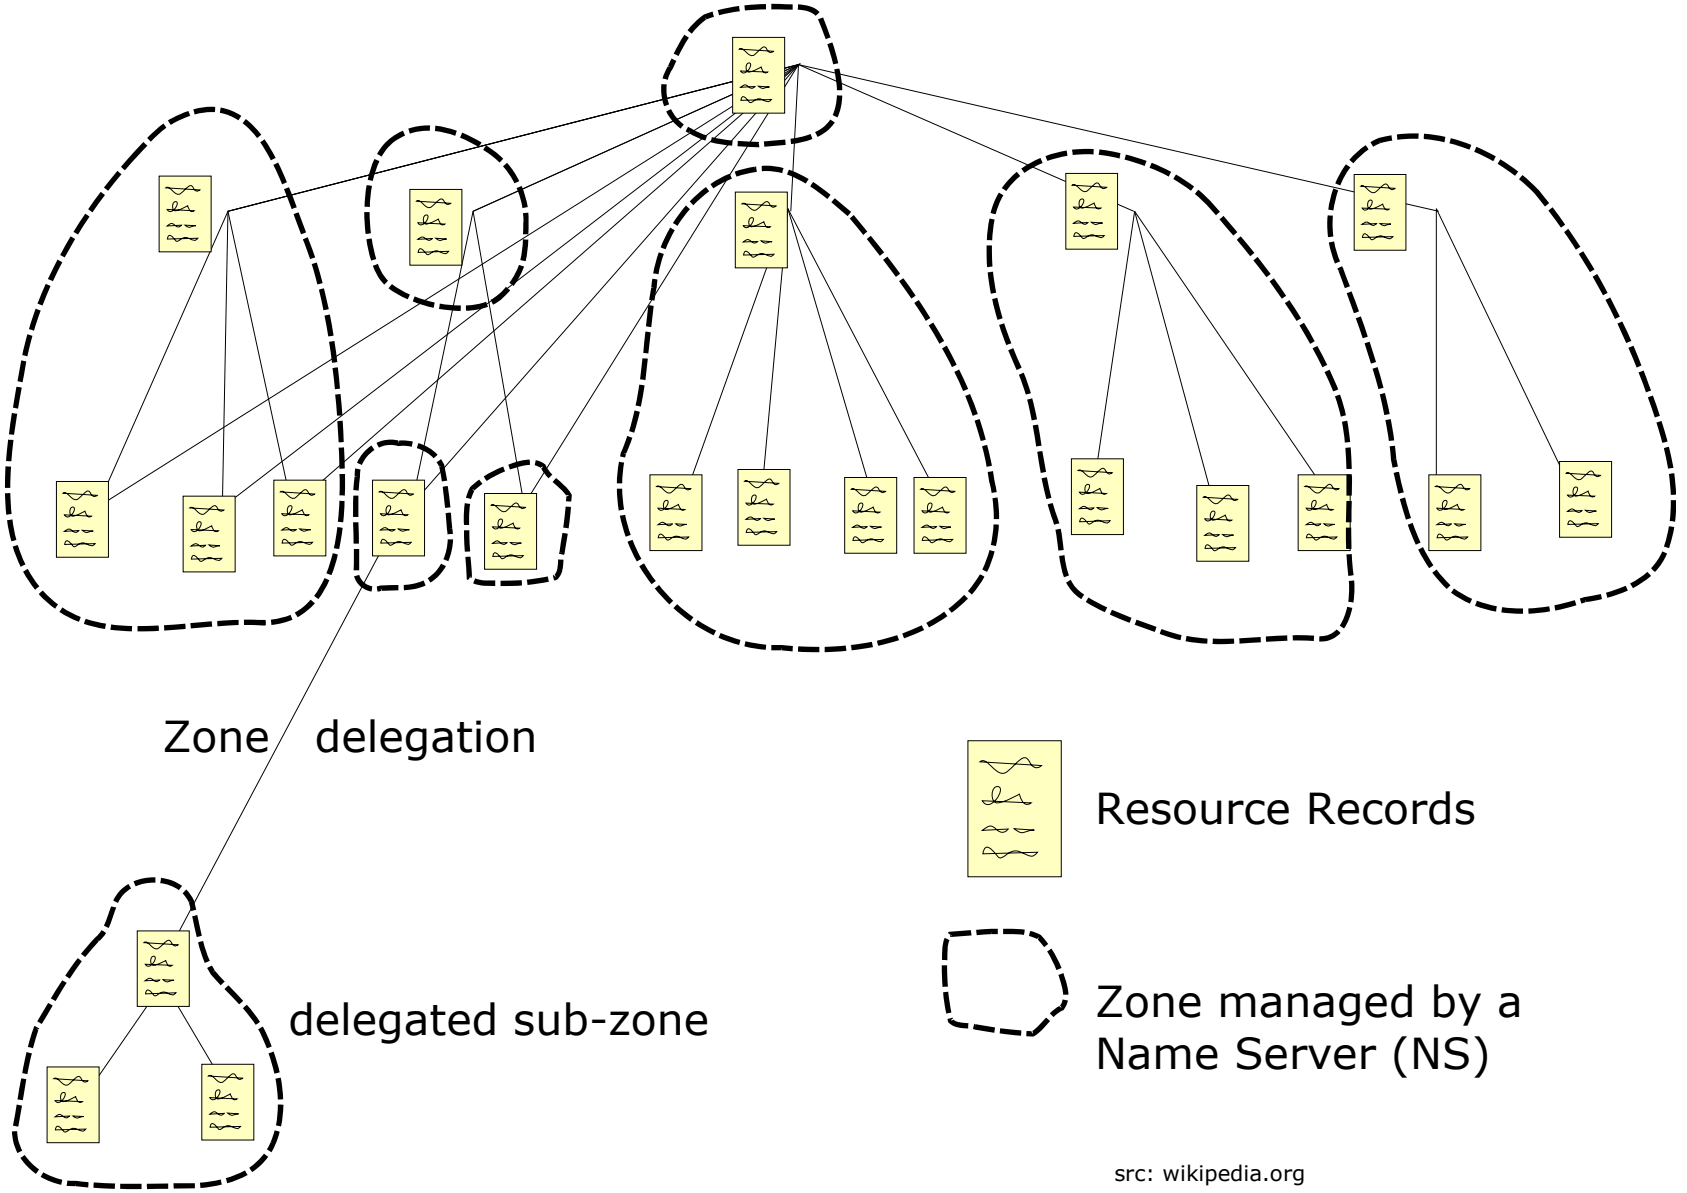

src: wikipedia.org

Tree structure of authorities (cont'd)

- Zone delegation is delegation of *trust*
- Example: the .se Name Server (NS) trusts the .kth.se NS to resolve *.kth.se
- .se has no say over .kth.se other than the delegation itself

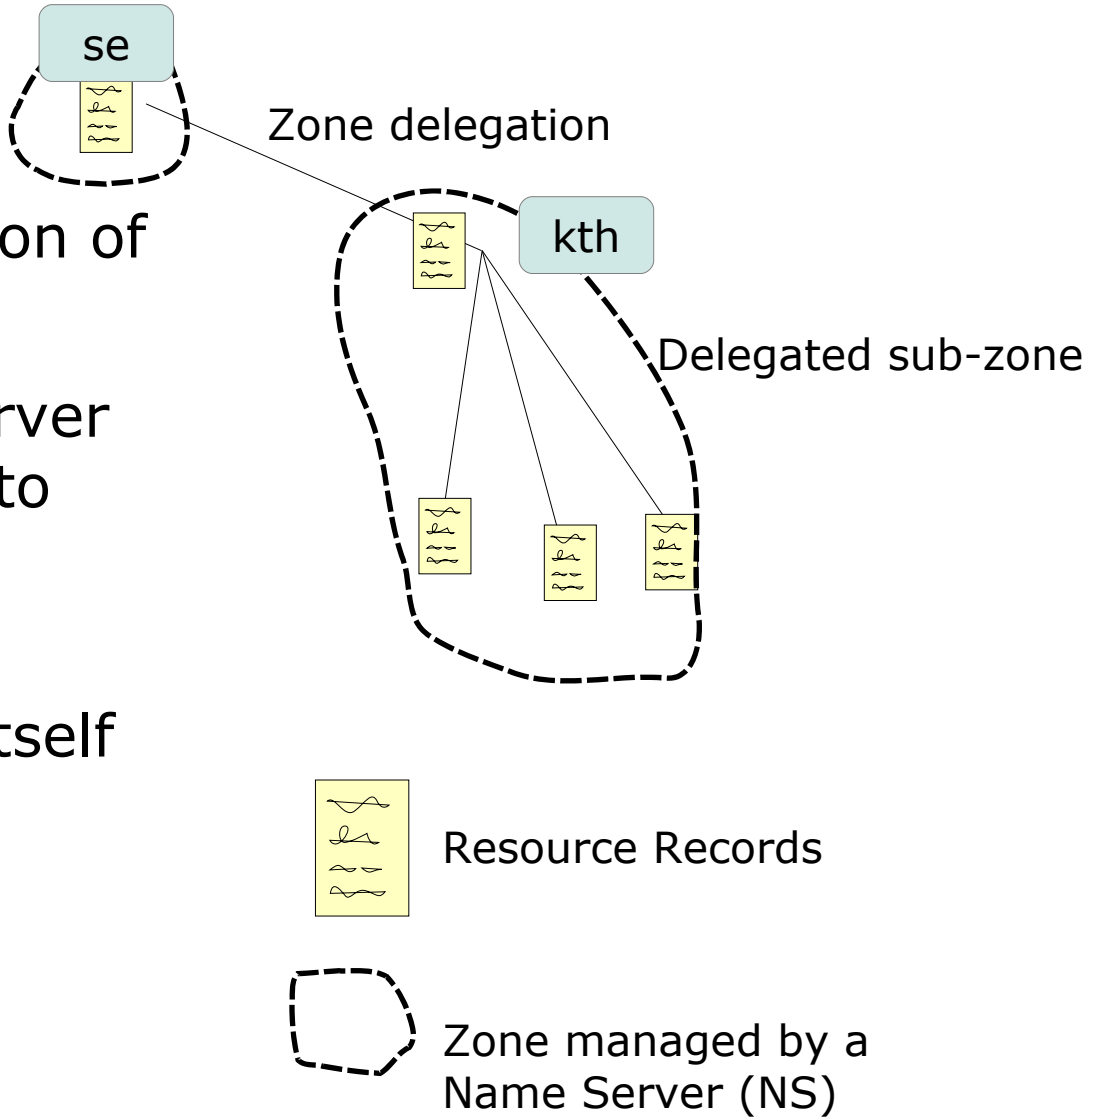

src: wikipedia.org

ROYAL INSTITUTE OF TECHNOLOGY

Recursive requests

- Clients/end-hosts get the answer from the nearest NS possible

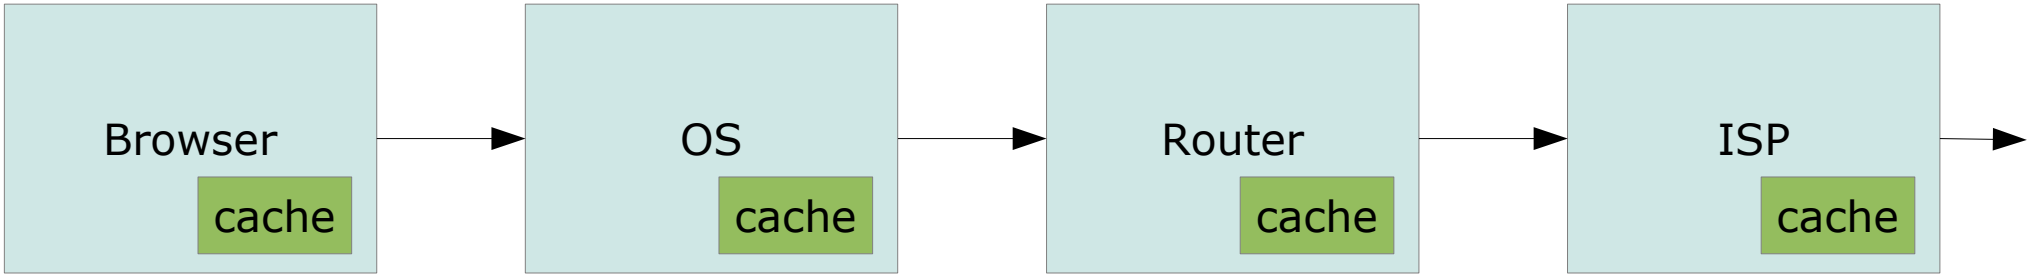

src: after wikipedia.org

Resolution example

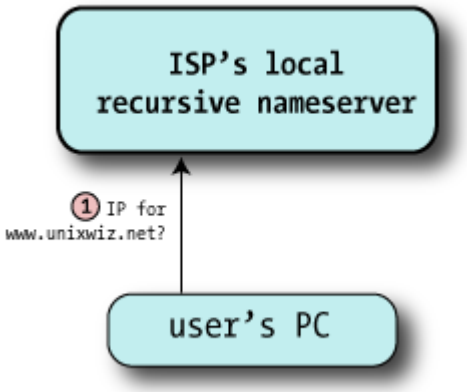

src: Steve Friedl, unixwiz.net

Resolution example (cont'd)

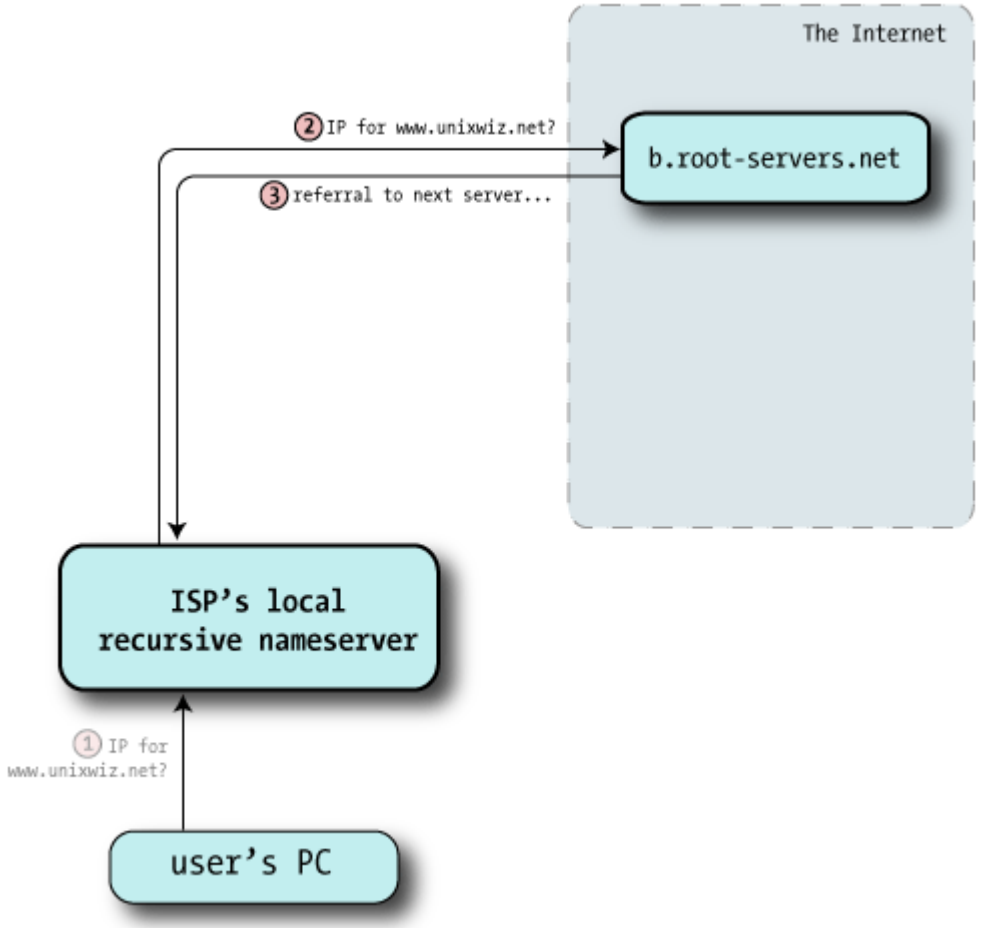

src: Steve Friedl, unixwiz.net

Resolution example (cont'd)

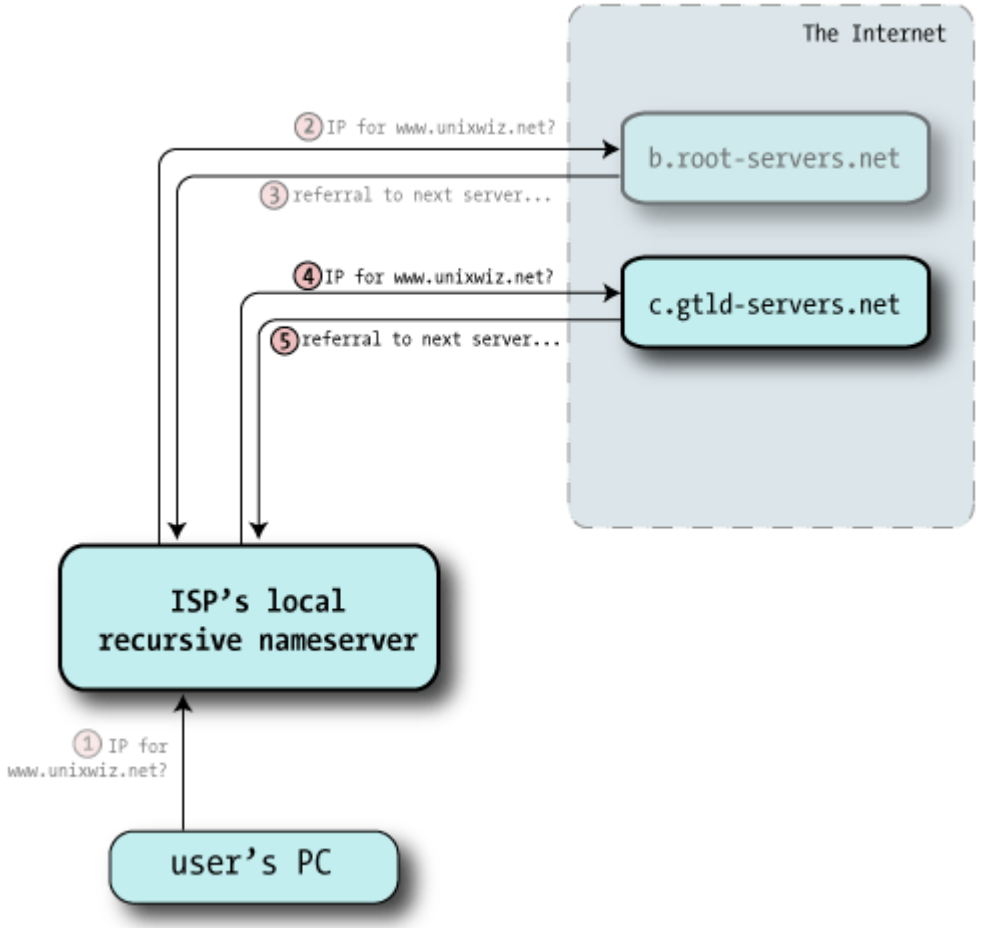

src: Steve Friedl, unixwiz.net

Resolution example (cont'd)

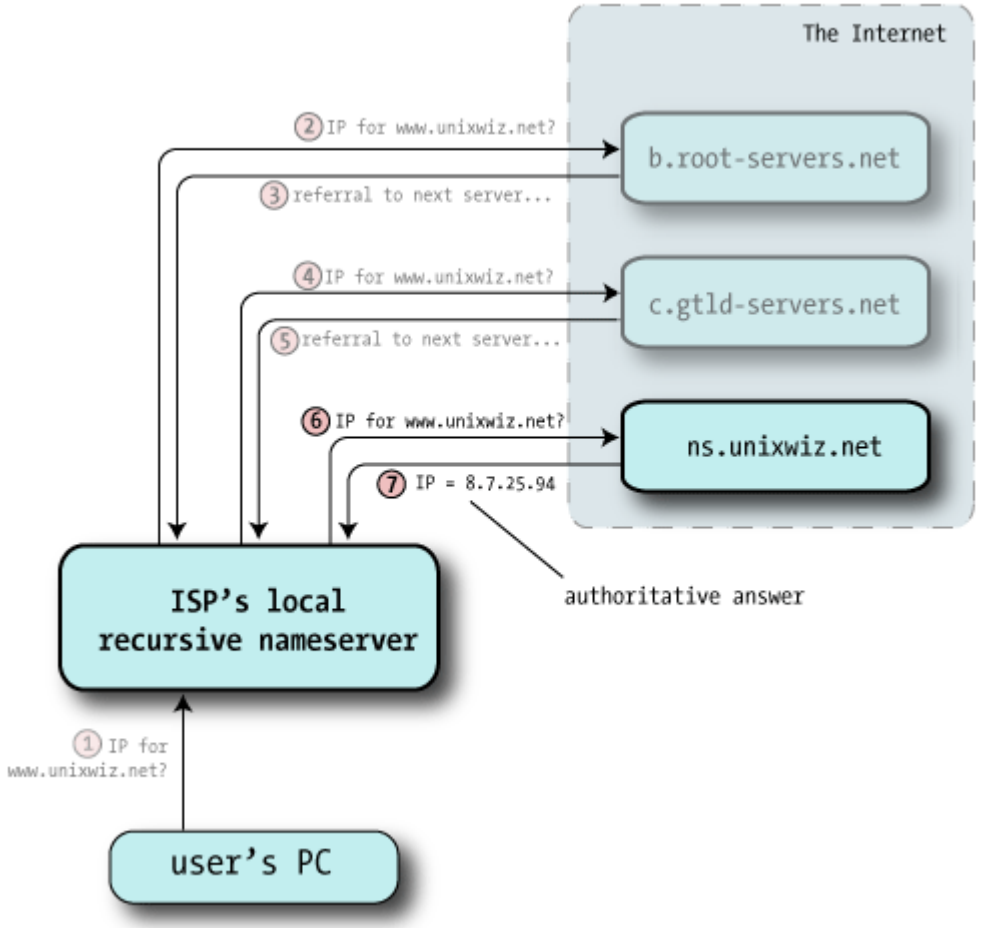

src: Steve Friedl, unixwiz.net

Resolution example (cont'd)

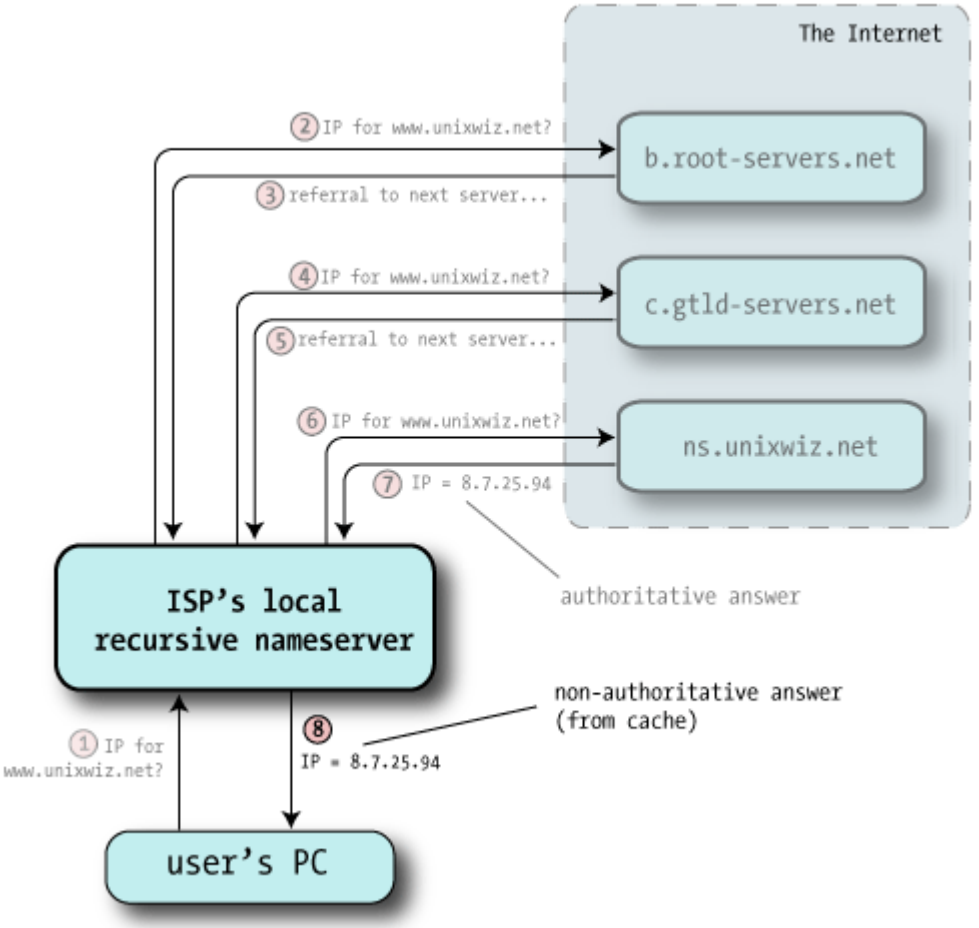

src: Steve Friedl, unixwiz.net

DNS query

- Query ID is unique per query and links it to the response
- 16 bits of randomizable data

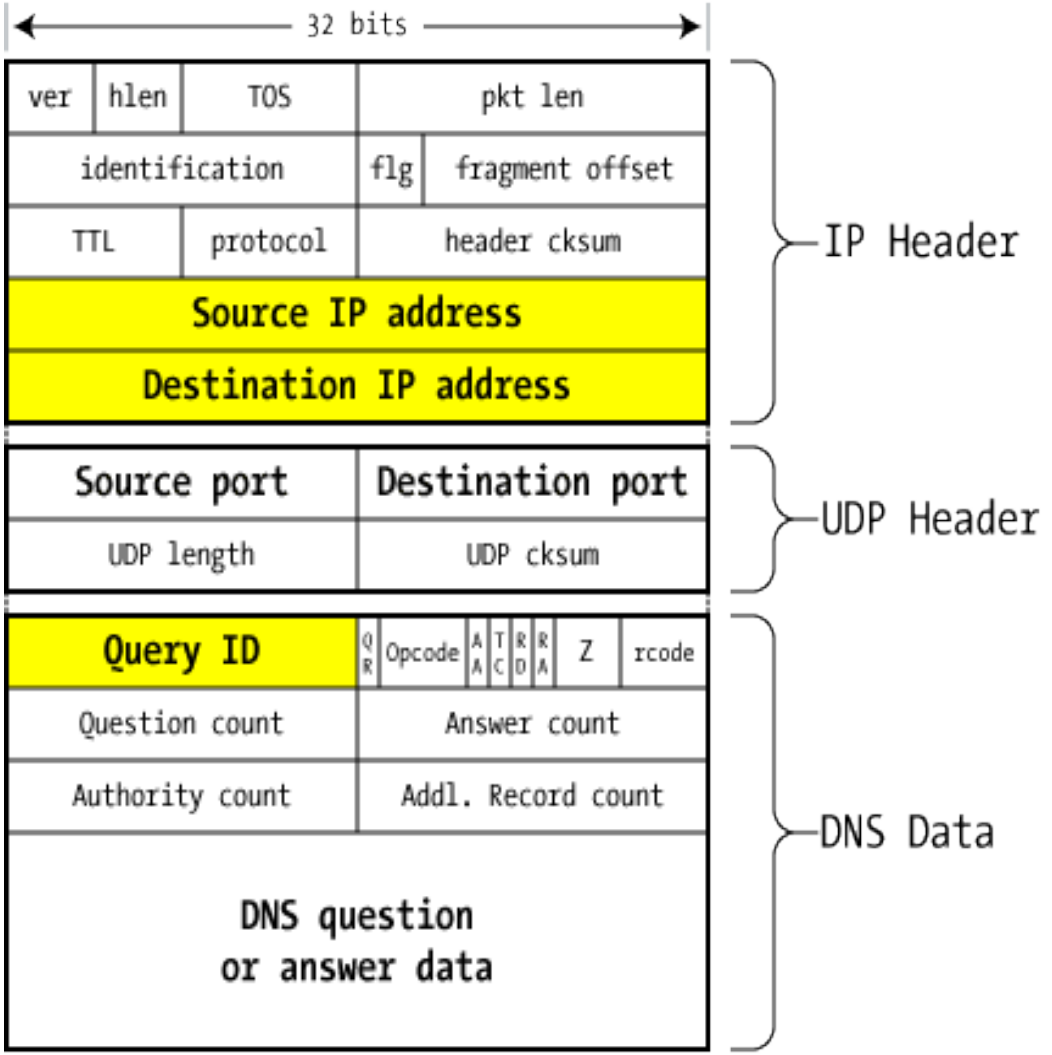

src: Steve Friedl, unixwiz.net

Query example

- DNS query to c.gtld-servers.net
- Gets QID 43561

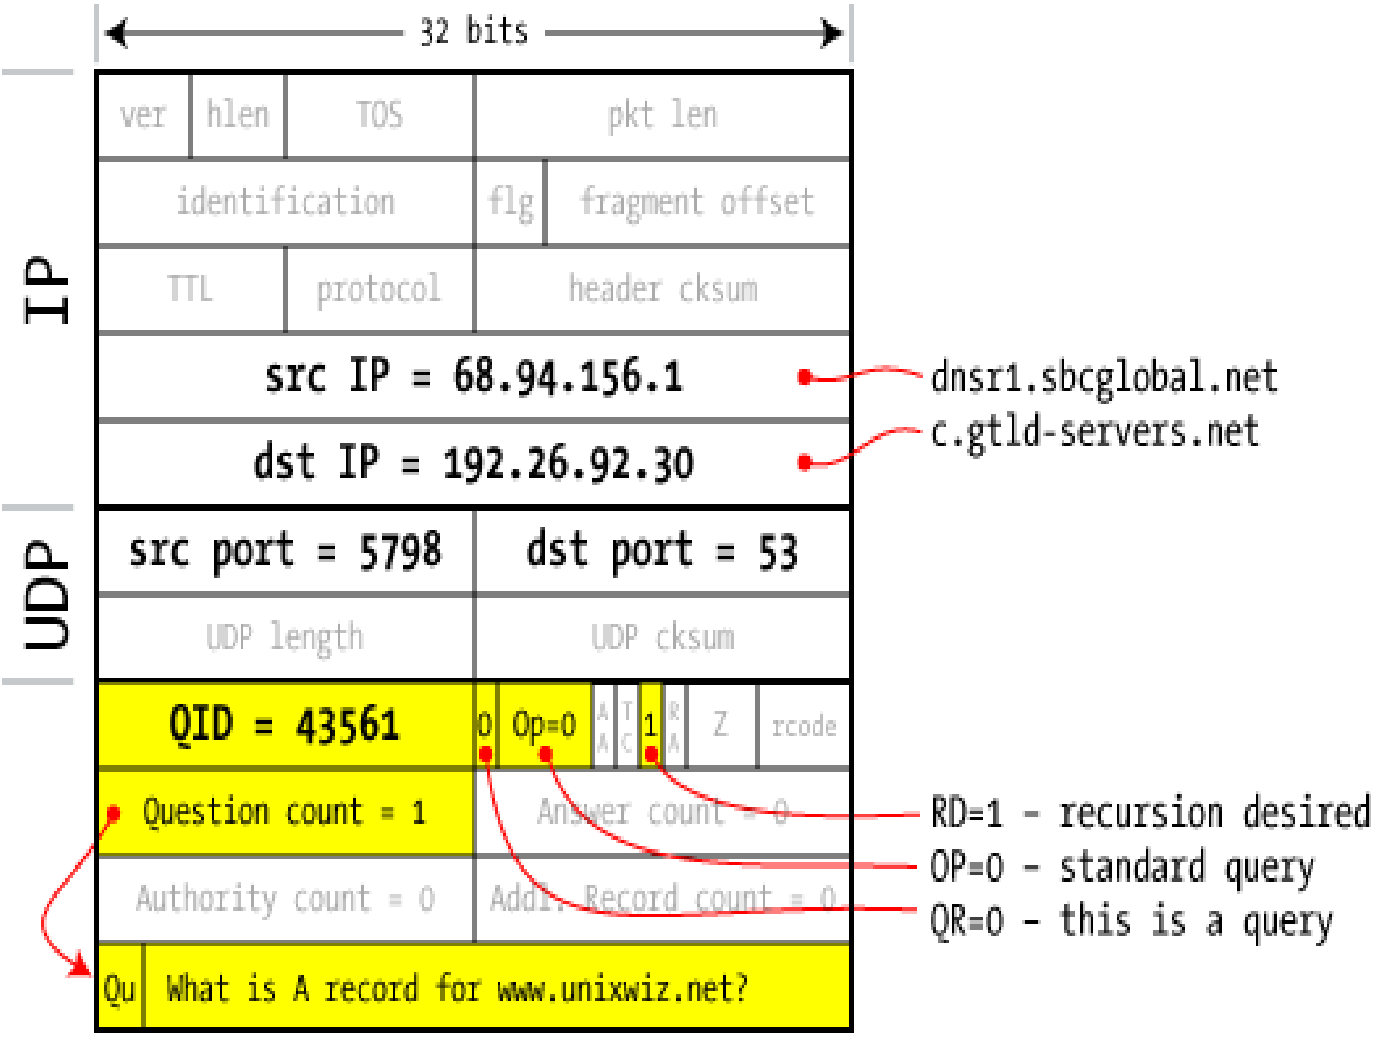

src: Steve Friedl, unixwiz.net

Query example (cont'd)

- DNS response from c.gtld-servers.net
- Identified by QID
- Unknown address: Response with next NS (name and IP)
- DNS glue records
  - IP addresses of NSs within the queried domain
  - What would happen without them?

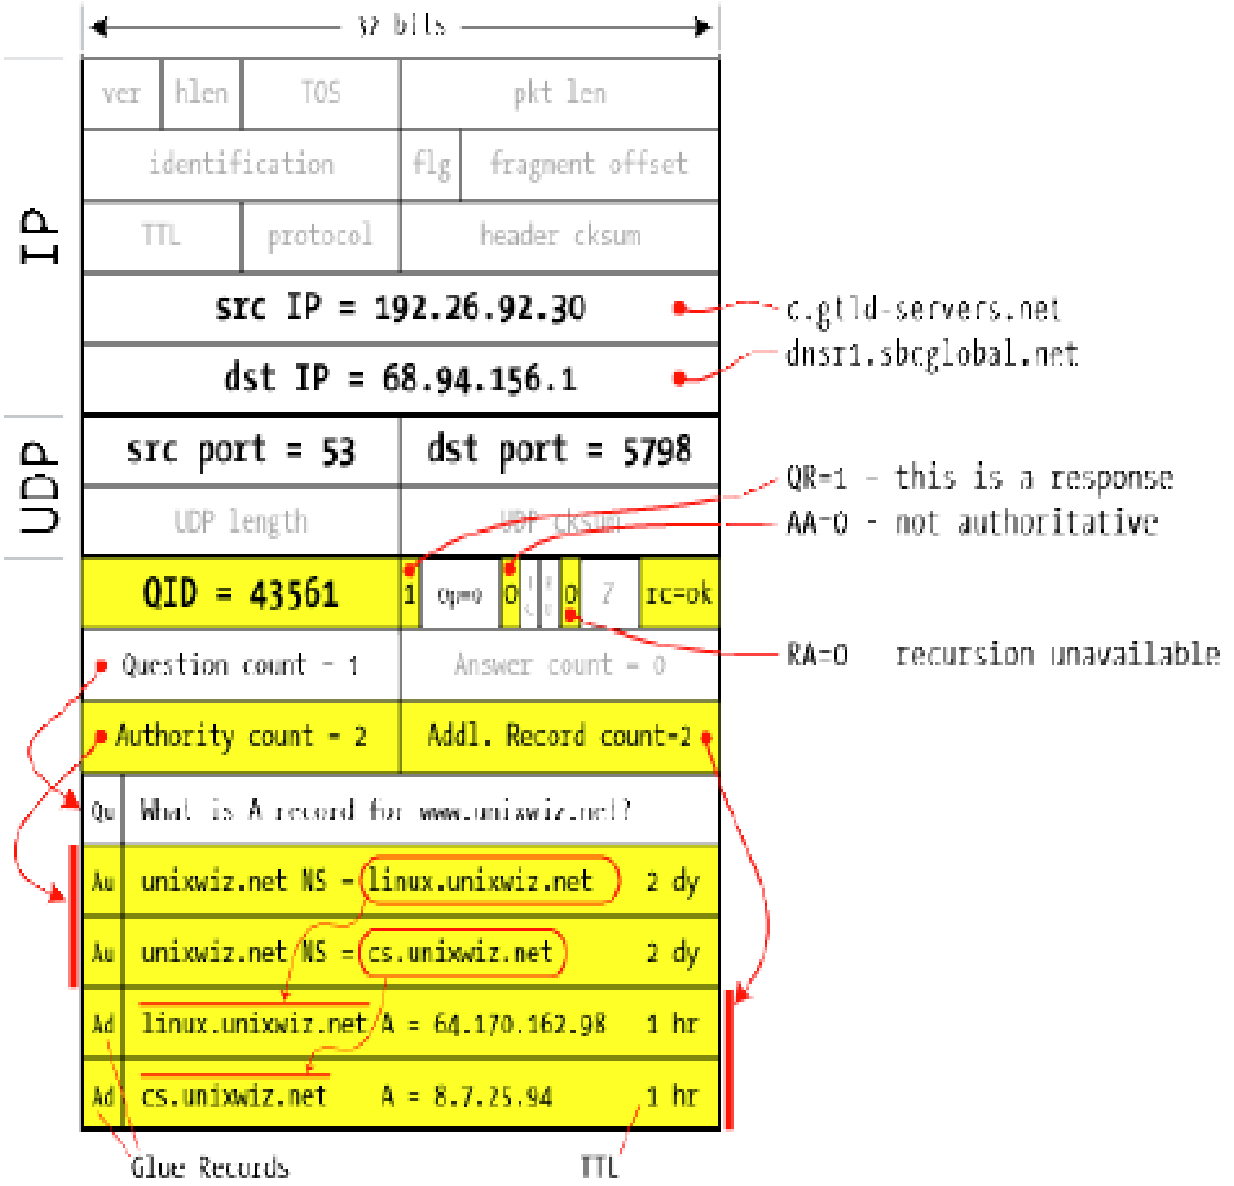

src: Steve Friedl, unixwiz.net

Query example (cont'd)

- DNS query to linux.unixwiz.net
- New QID
- Same sort of request

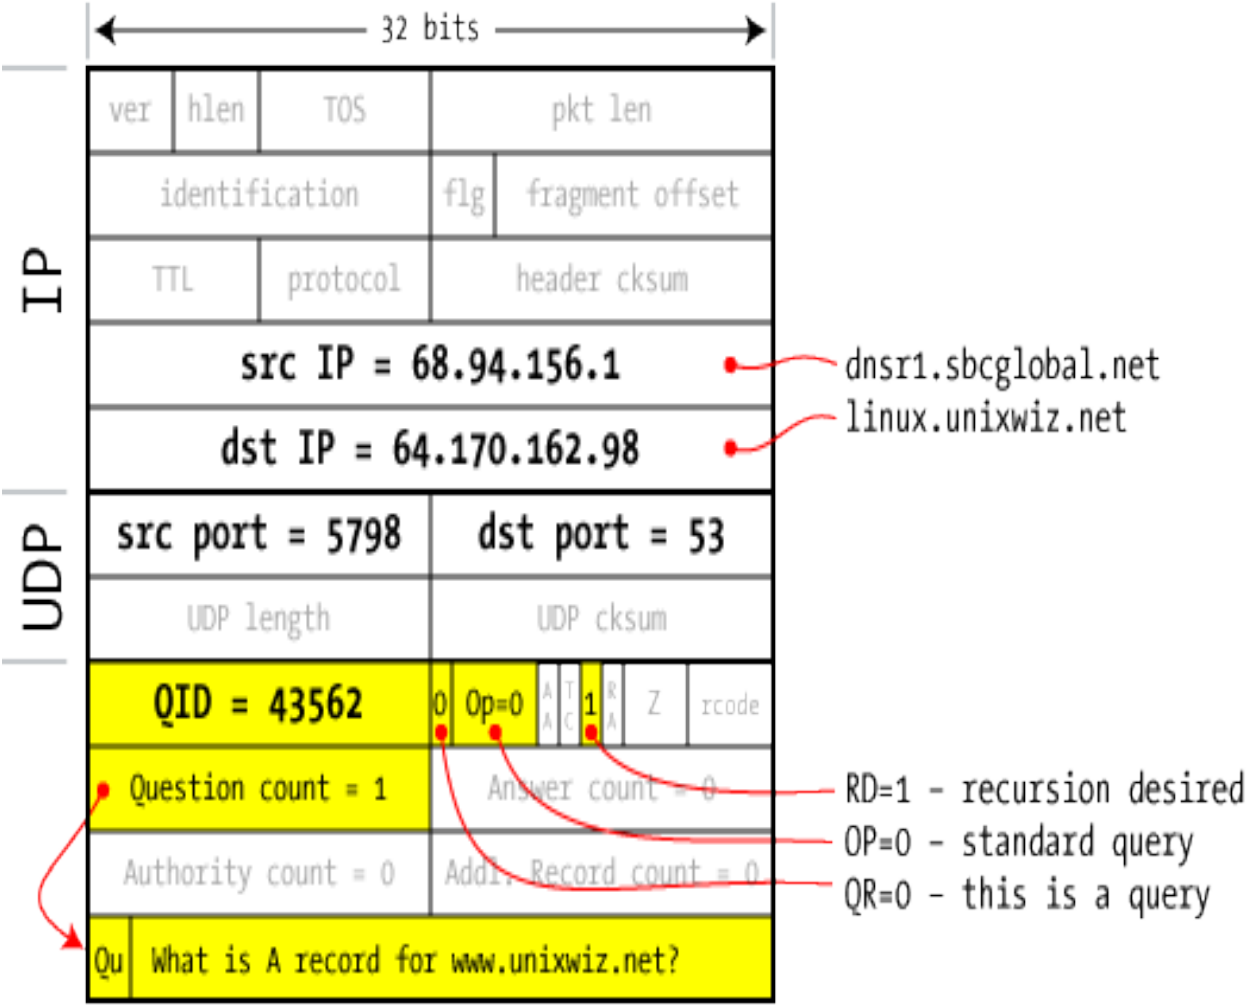

src: Steve Friedl, unixwiz.net

Query example (cont'd)

- DNS authoritative response from linux.unixwiz.net
- Linked by QID
- We got our answer

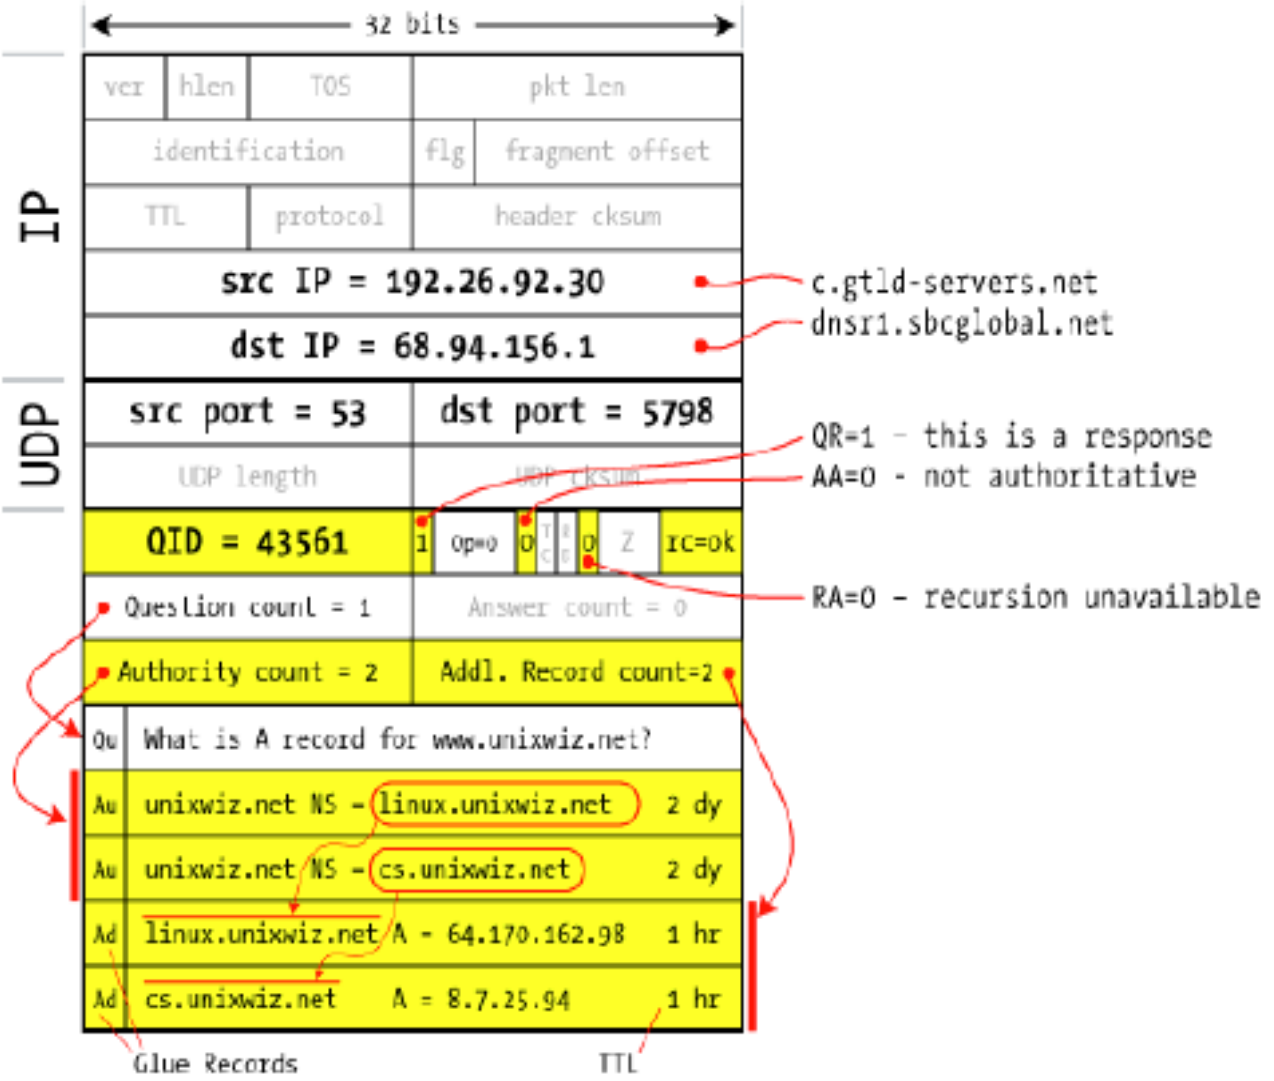

src: Steve Friedl, unixwiz.net

Checking the response

- Same UDP port we sent it from

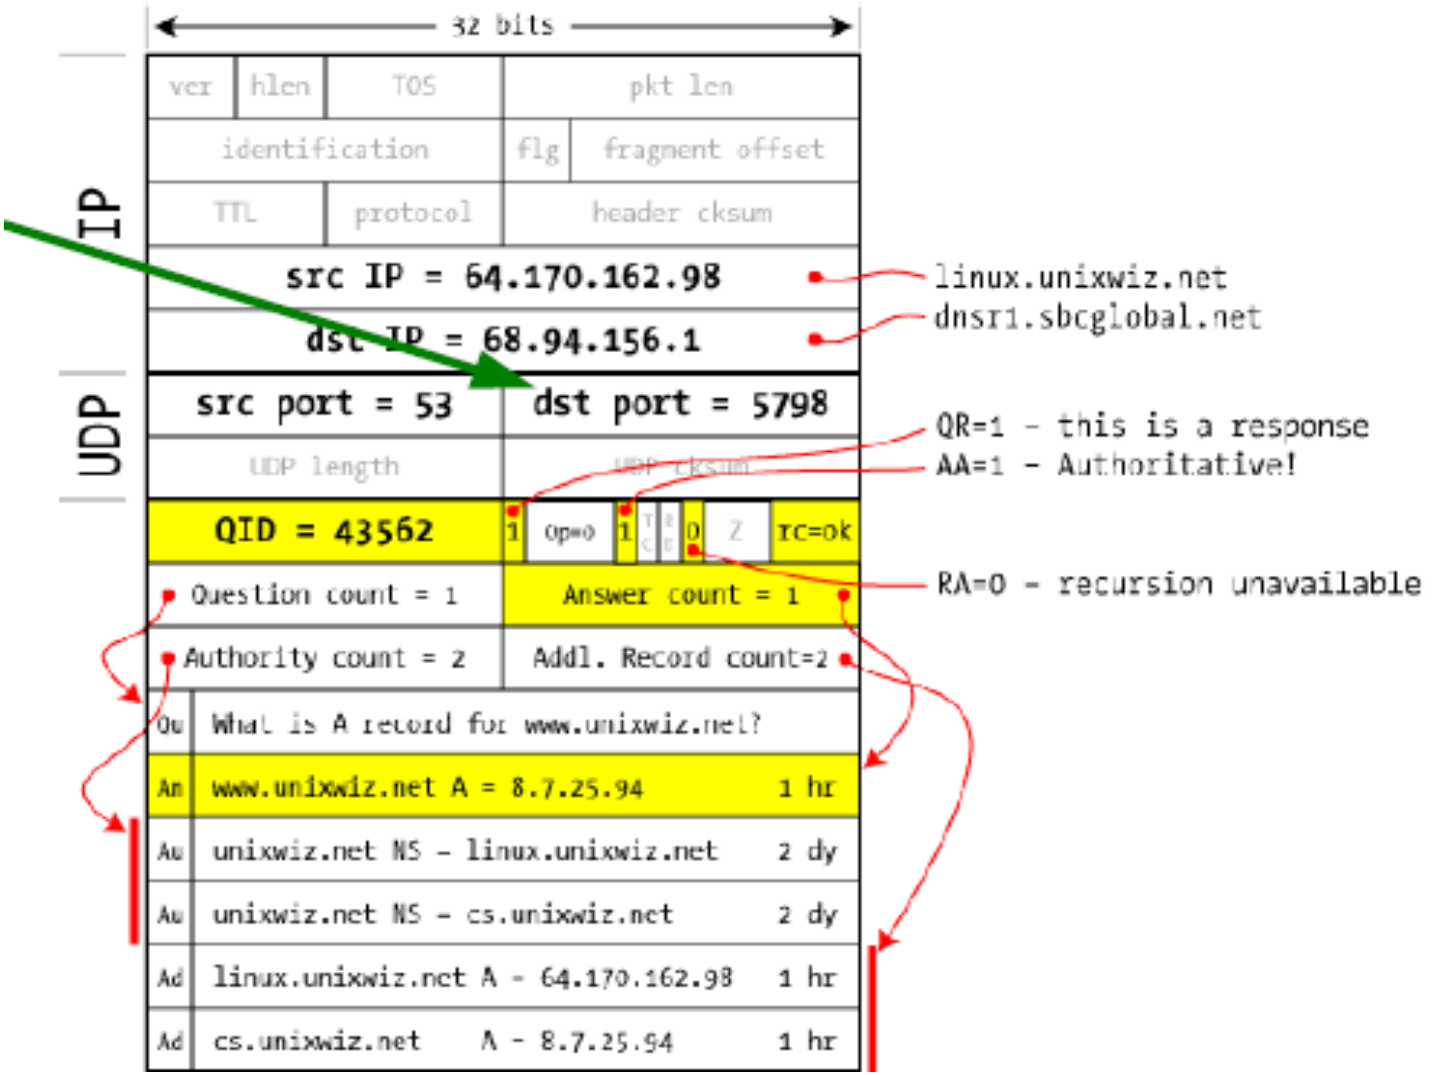

src: Steve Friedl, unixwiz.net

Checking the response (cont'd)

- The Query ID matches the pending query

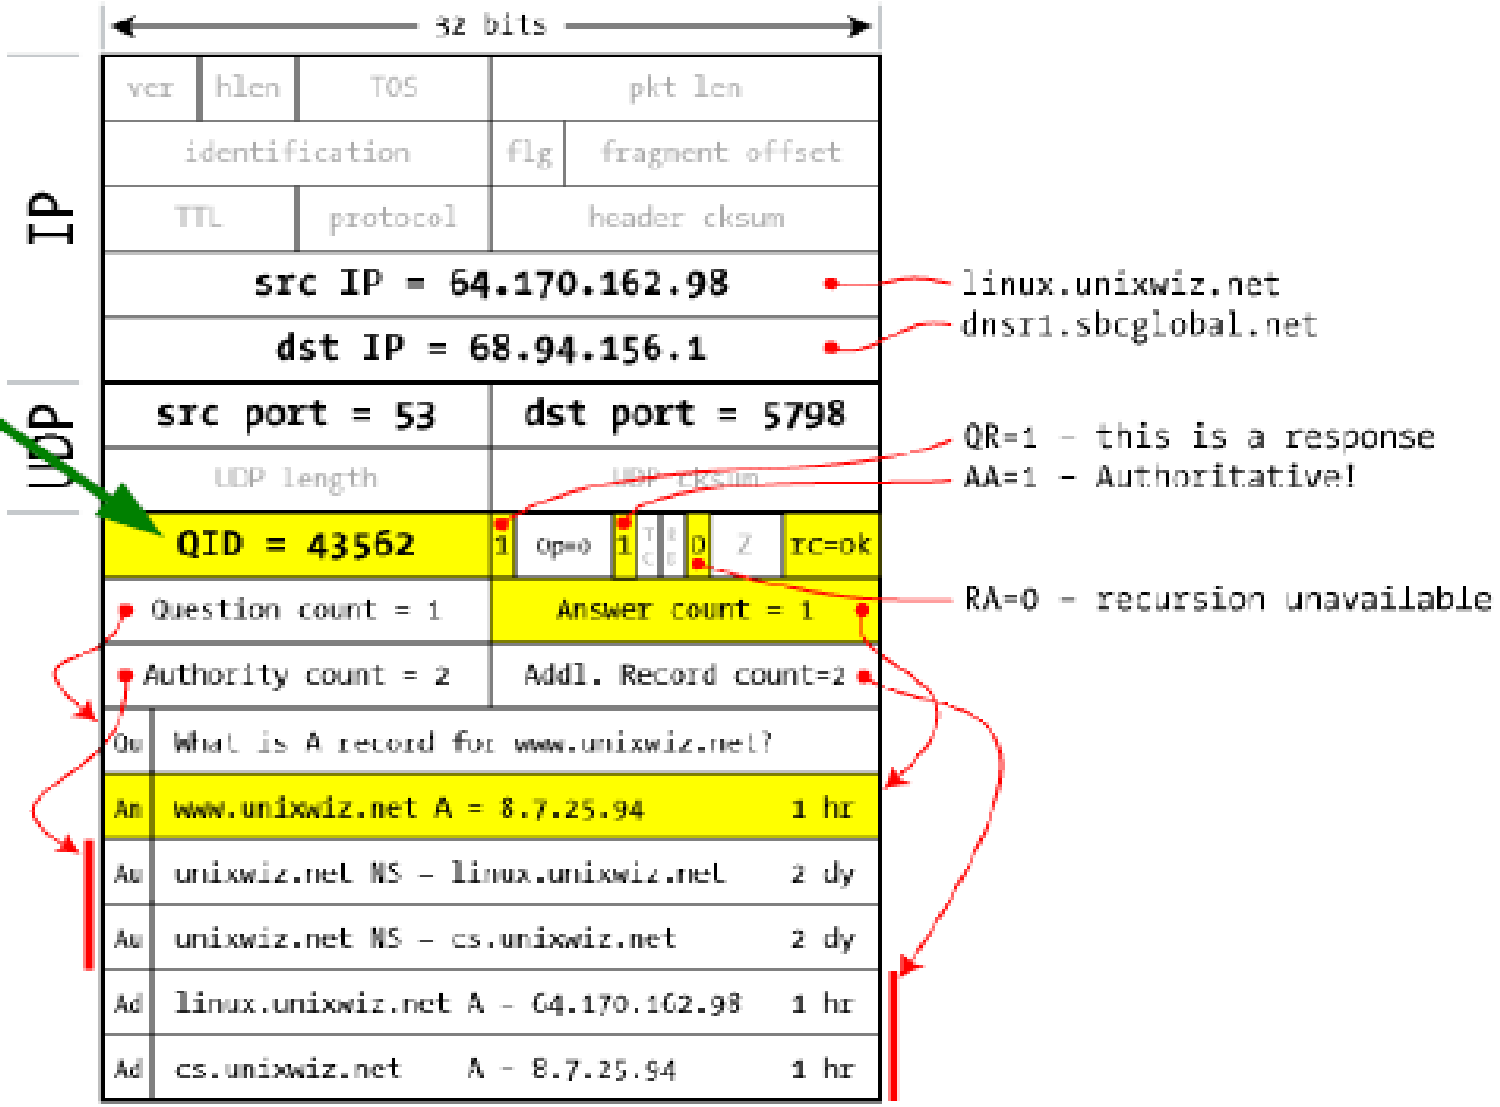

src: Steve Friedl, unixwiz.net

Checking the response (cont'd)

- The Question section is a duplicate

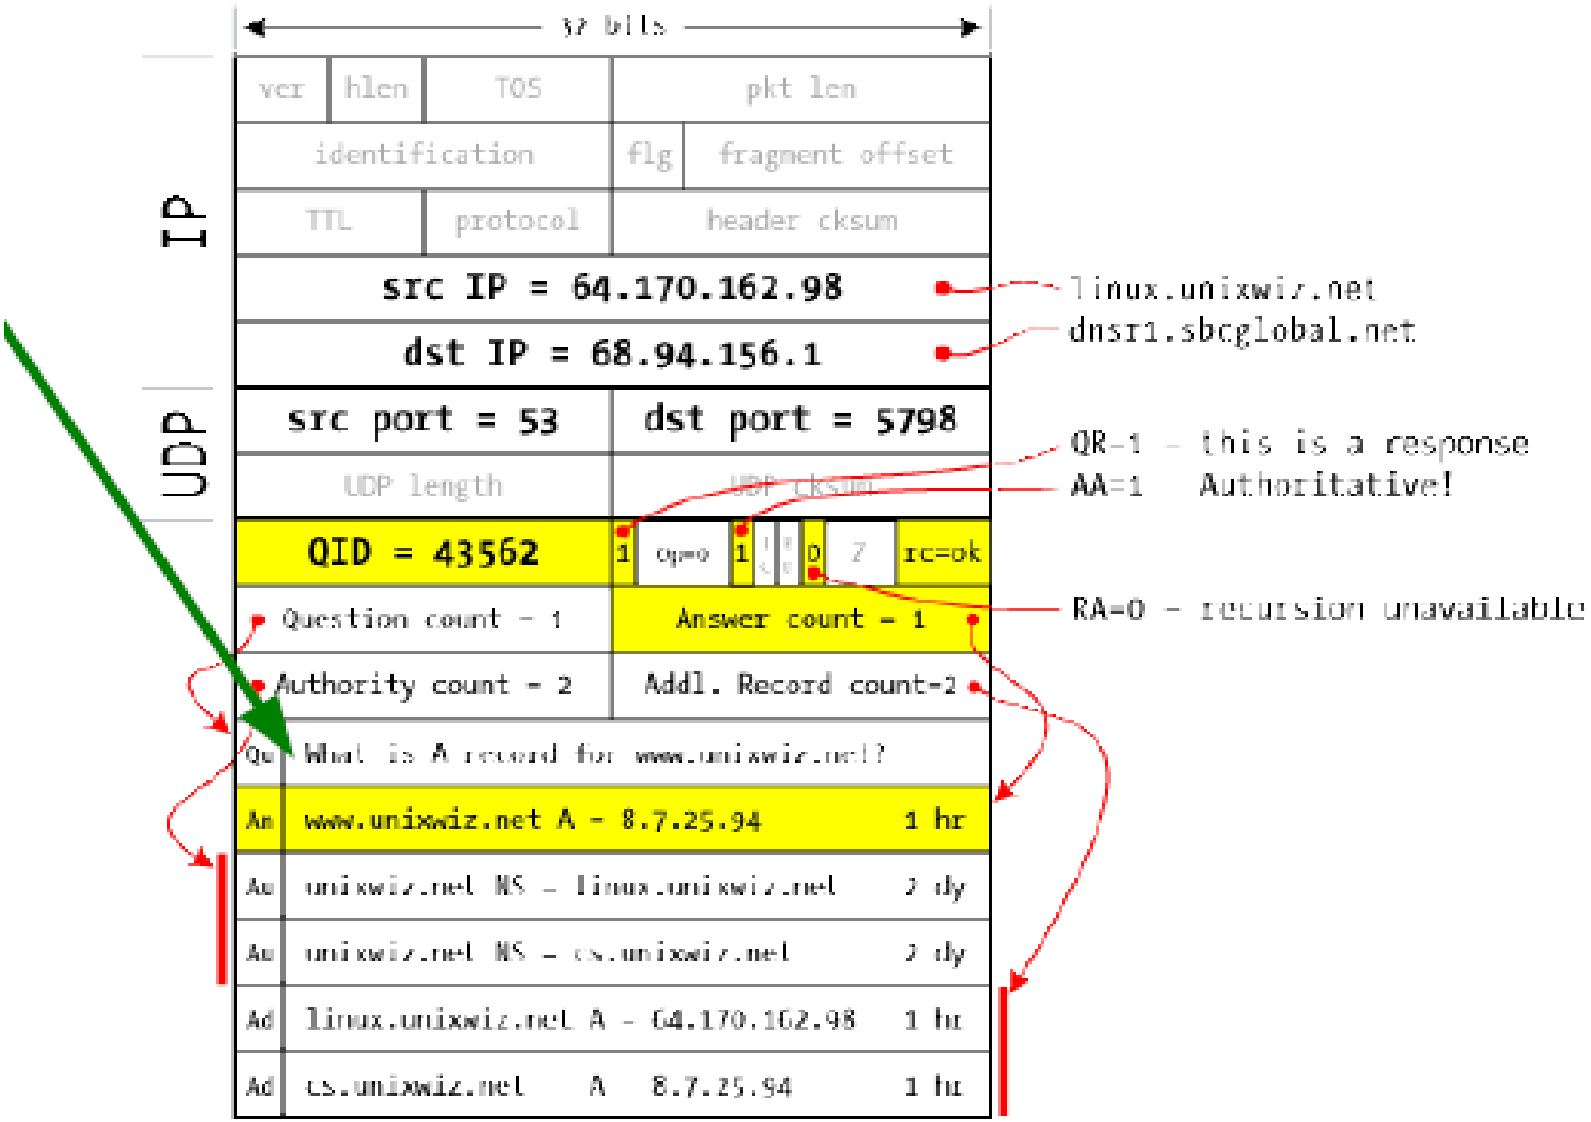

src: Steve Friedl, unixwiz.net

Checking the response (cont'd)

- Response is in the same domain as the query ("bailiwick checking")

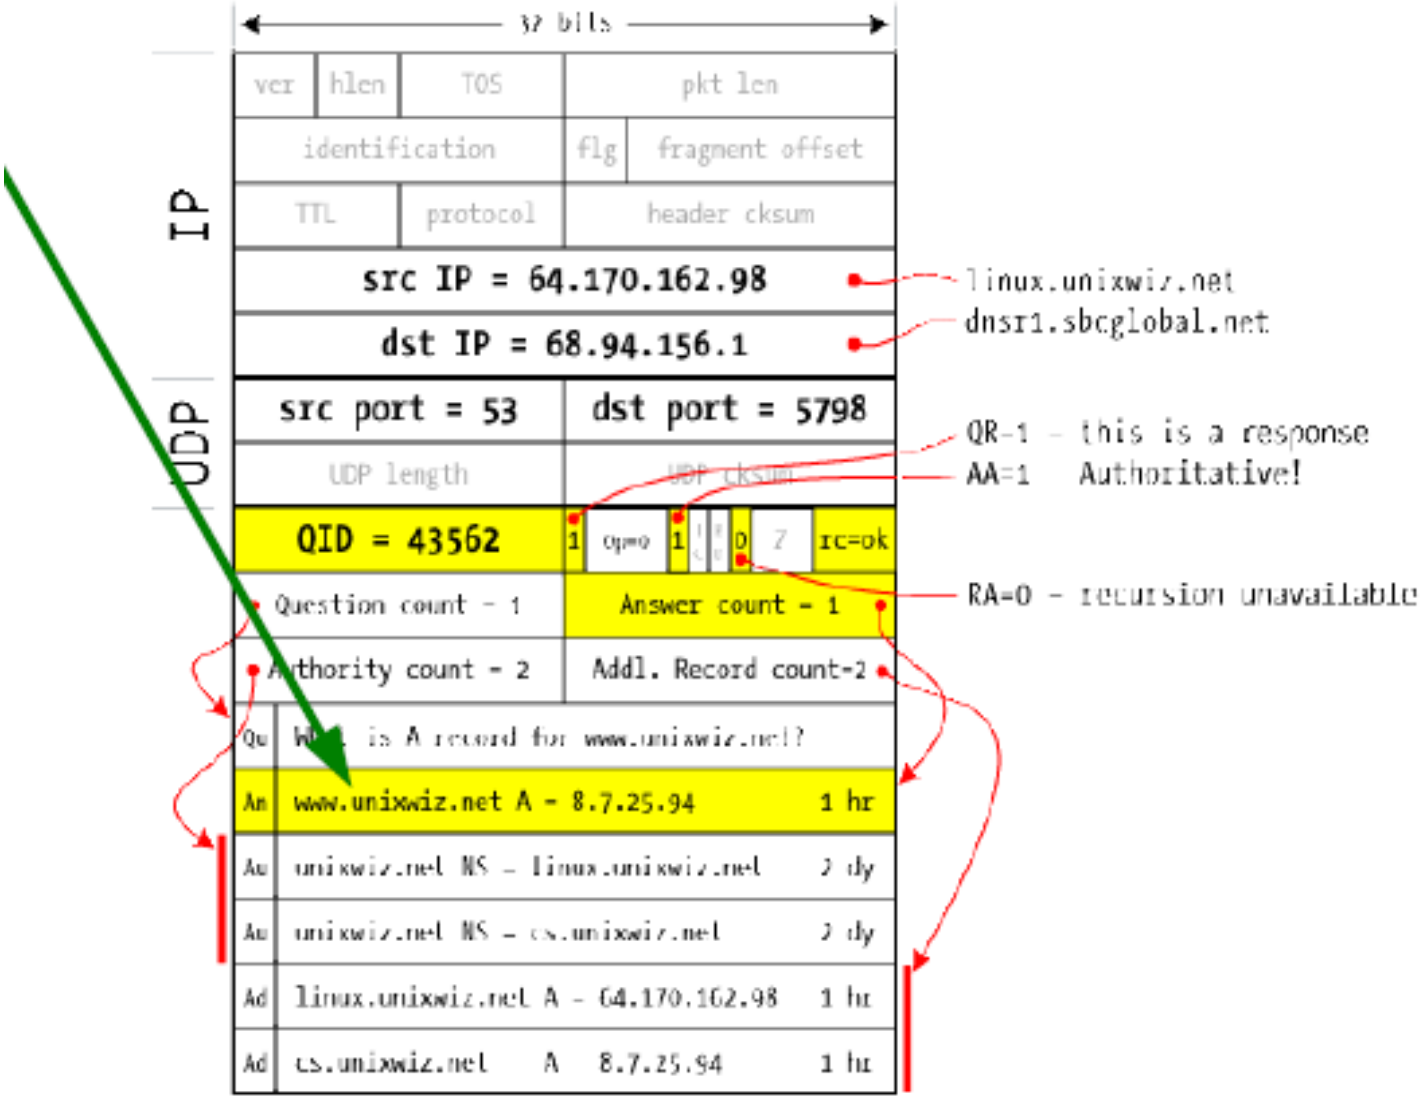

src: Steve Friedl, unixwiz.net

Time To Live (TTL)

- A DNS answer contains a TTL describing how long to keep the record
- Answers are kept in caches
- In a way, the responder manages the cache

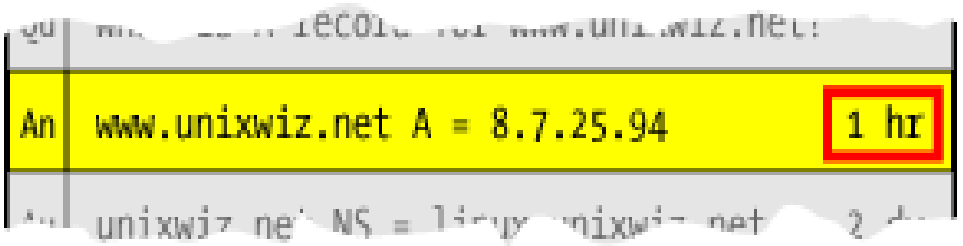

An	www.unixwiz.net A = 8.7.25.94	1 hr

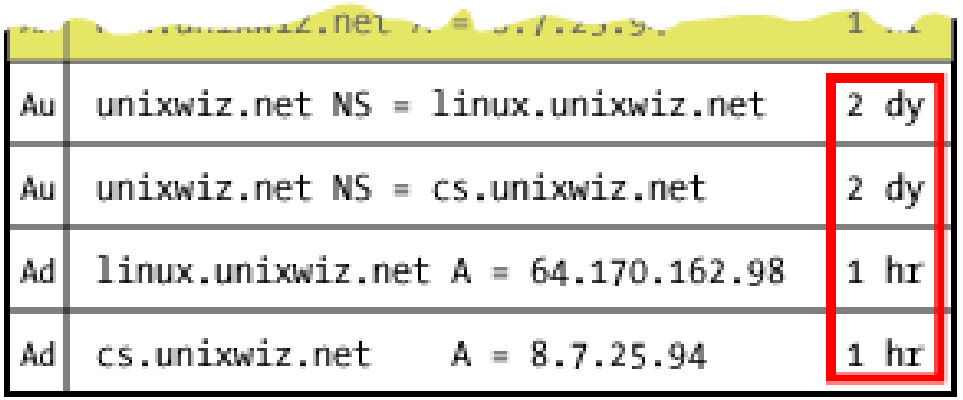

Au	unixwiz.net NS = linux.unixwiz.net	2 dy
Au	unixwiz.net NS = cs.unixwiz.net	2 dy
Ad	linux.unixwiz.net A = 64.170.162.98	1 hr
Ad	cs.unixwiz.net A = 8.7.25.94	1 hr

How can we inject malicious data?

src: Steve Friedl, unixwiz.net

ROYAL INSTITUTE OF TECHNOLOGY

DNS Attacks

- How can we abuse DNS?
- *Examples:*
  - Man-in-the-Middle attacks
  - Kaminsky cache poisoning
  - DNS rebinding
- *Recommended reading:*
  - RFC3833: Threat Analysis of the Domain Name System (DNS)

ROYAL INSTITUTE OF TECHNOLOGY

DNS Attacks (cont'd)

- Why would we want to attack DNS?
  - To pretend to be someone else
  - To redirect users to where we want them to
  - Fun and profit

Man-in-the-Middle

- The adversary can change the response, drop it, or create its own in an arbitrary manner
- Sometimes used by governments, asking ISPs to
  - Provide “erroneous” responses
  - Drop queries for specific sites

ROYAL INSTITUTE OF TECHNOLOGY

Government-in-the-Middle

- The Pirate Bay, along with several torrent sites, has been DNS-blocked in several countries
  - Is this how DNS was meant to be used?

Government-in-the-Middle (cont'd)

- In 2007 koreabonsai.com was blocked by Swedish ISPs, since the police listed it as child pornography
- Such lists mostly block their intended targets, but is there due process?

Government-in-the-Middle (cont'd)

- SOPA/PIPA in the USA contained DNS blocking sections, causing Wikipedia and others to blackout in protest

ROYAL INSTITUTE OF TECHNOLOGY

Kaminsky cache poisoning

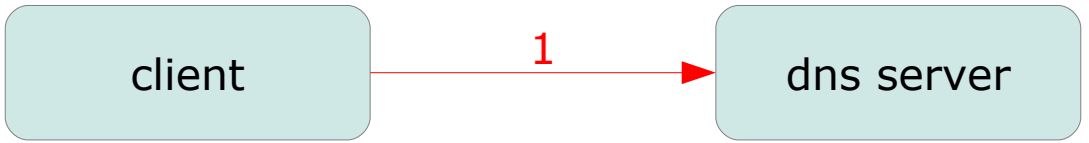

- Let client query (1) for random.trusted.com

Kaminsky cache poisoning (cont'd)

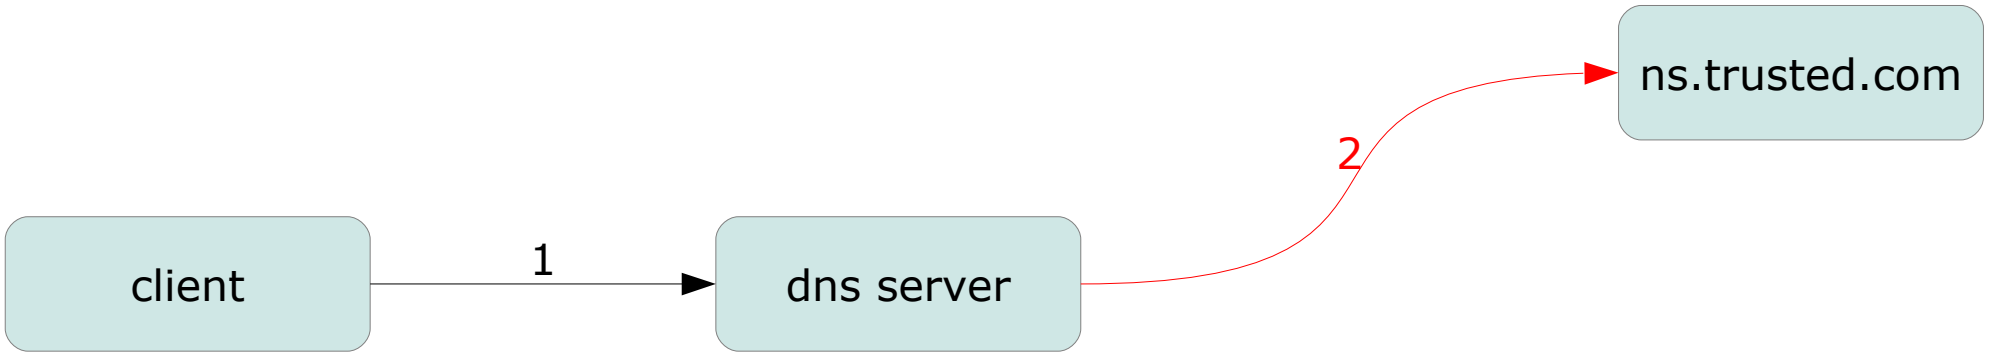

- Let client query for random.trusted.com (1)
- The local dns server will ask/look for ns.trusted.com (2)

Kaminsky cache poisoning (cont'd)

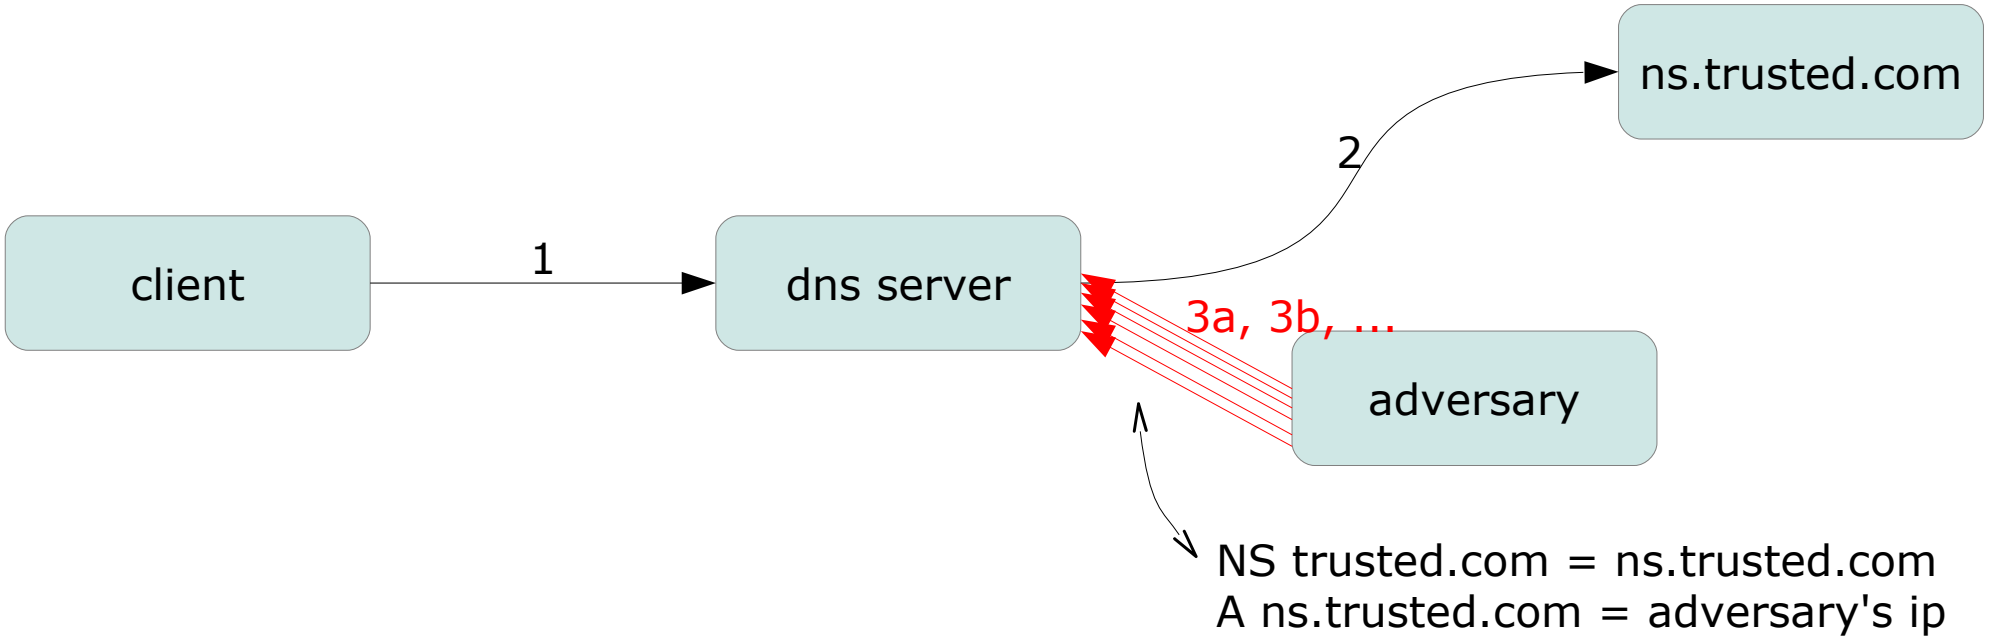

- Let client query for random.trusted.com (1)
- The local DNS server will ask/look for ns.trusted.com (2)
- The attacker sends multiple responses (3) with different QID. If any one (3) matches (2), the attacker will now "own" trusted.com

Kaminsky cache poisoning (cont'd)

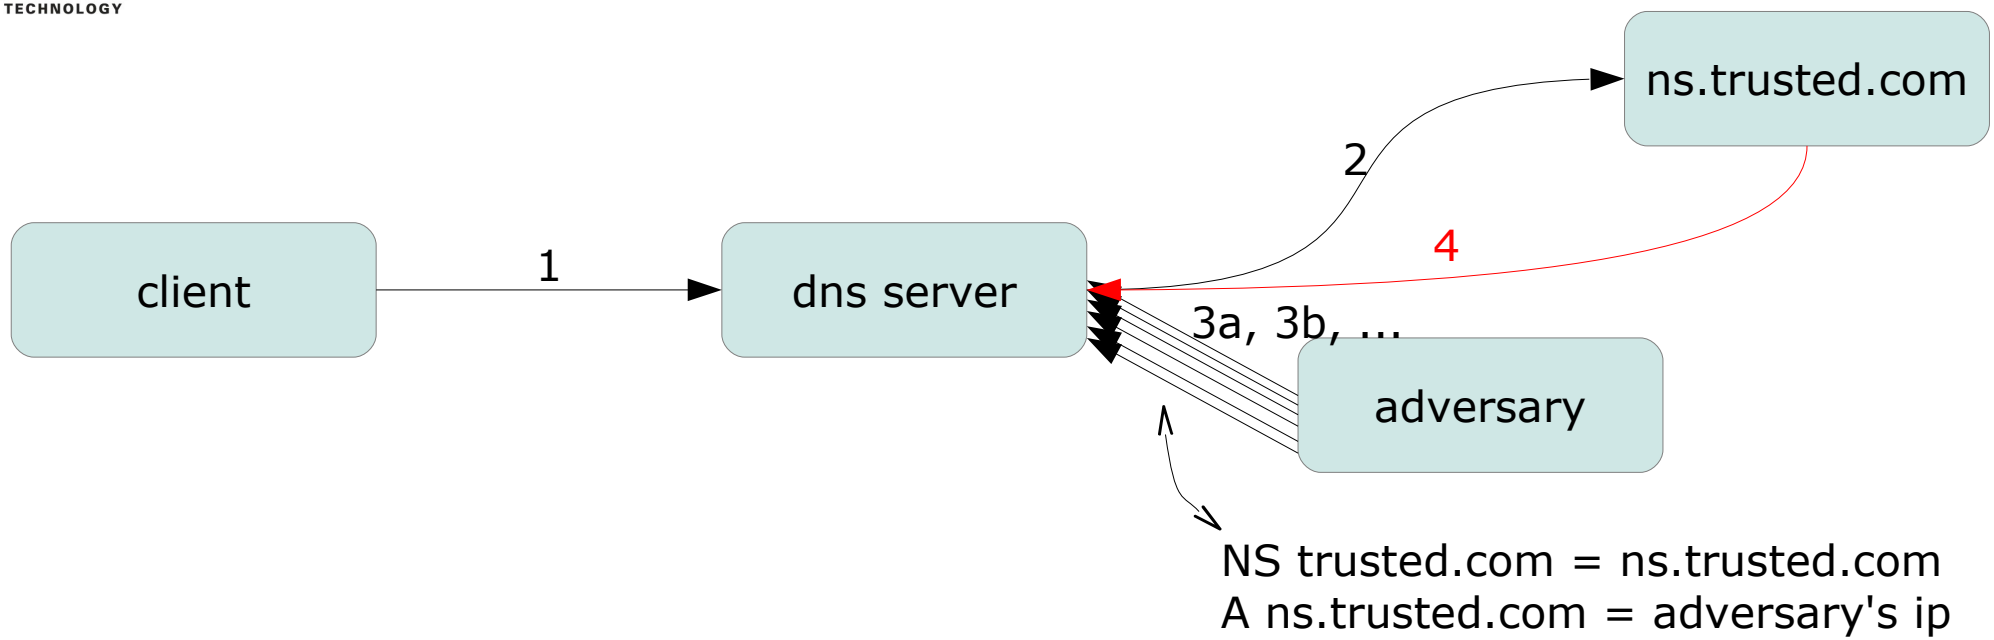

- Let client query for random.trusted.com (1)
- The local DNS server will ask/look for ns.trusted.com (2)
- The attacker sends multiple responses (3) with different QID. If any one (3) matches (2), the attacker will now "own" trusted.com
- Unless (4) arrives first

Kaminsky cache poisoning (cont'd)

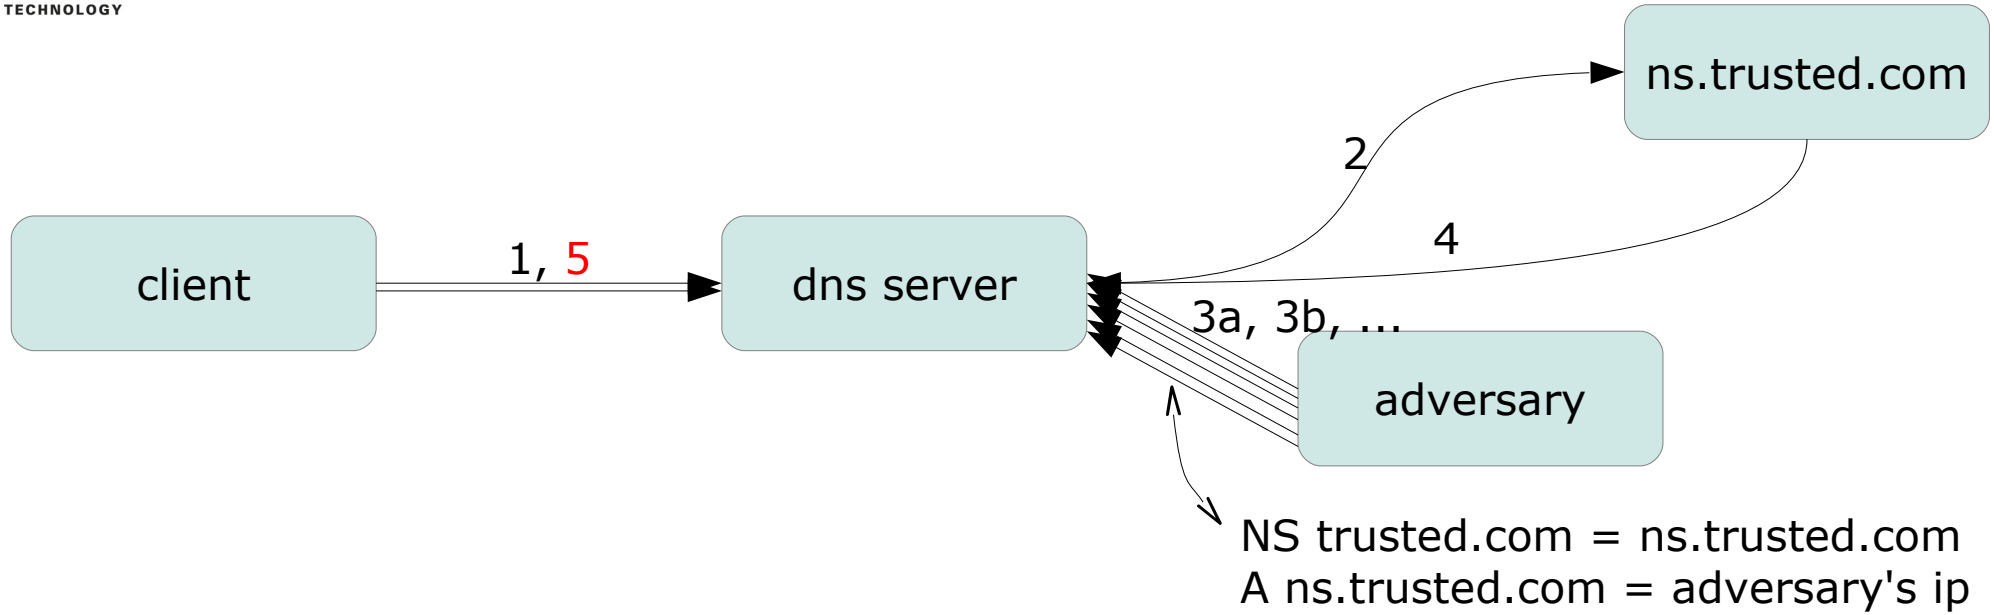

- ...
- Unless (4) arrives first
- If the attack fails, let the client query (5) for random2.trusted.com

Mitigation: Increase randomization

- Query ID is 16 bits
- Source port is 16 bits, and the DNS server can allocate a range of them, e.g., 11 bits
- 2^27 is much bigger than 2^16
- DNSSEC would solve this, but still has not been fully deployed

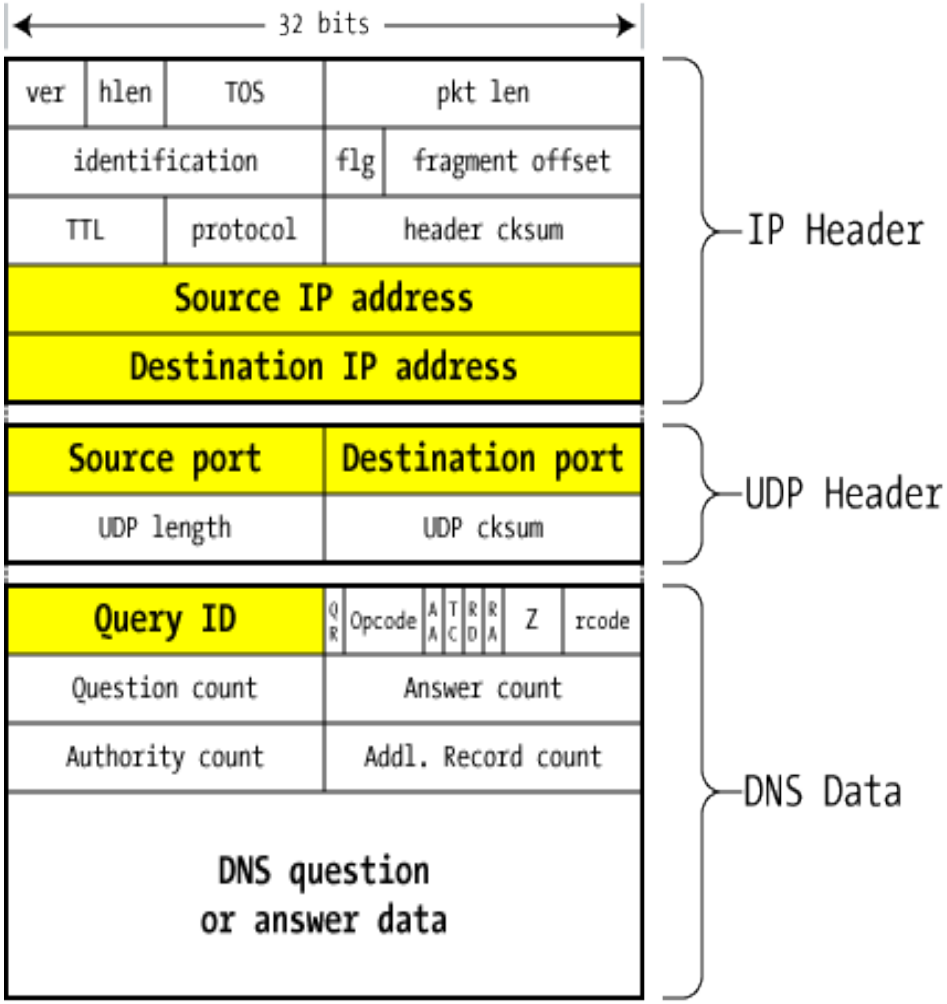

src: Steve Friedl, unixwiz.net

DNS Rebinding

- (1) is query for www.adversary

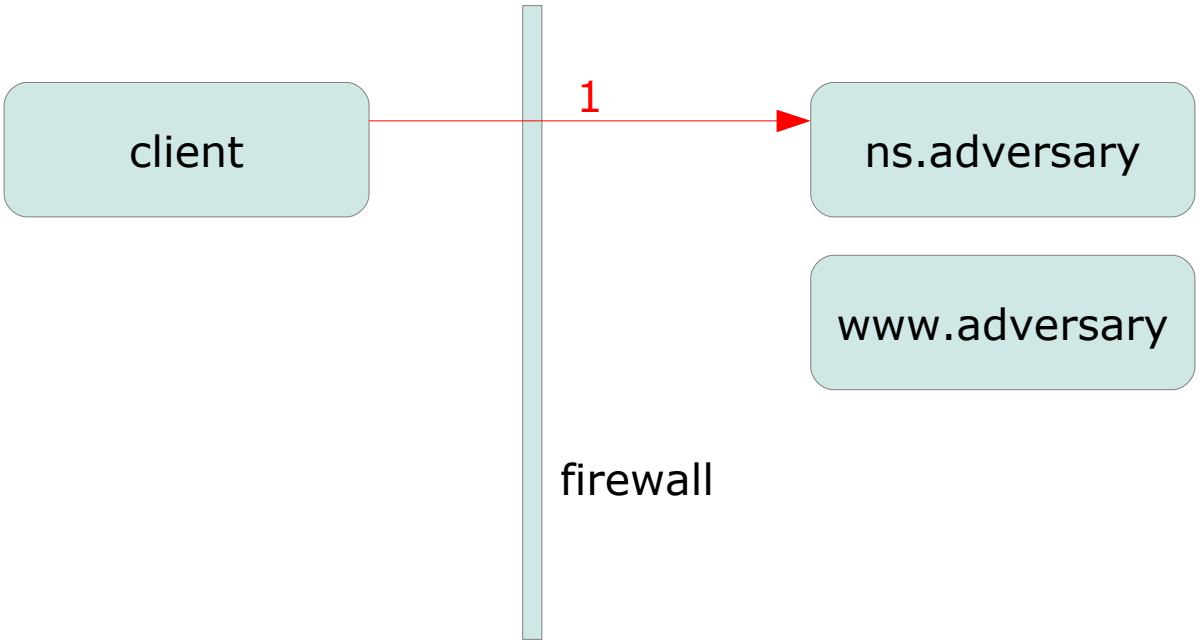

DNS Rebinding (cont'd)

- (1) is query for www.adversary
- (2) is a correct response pointing to (3), but with a short TTL

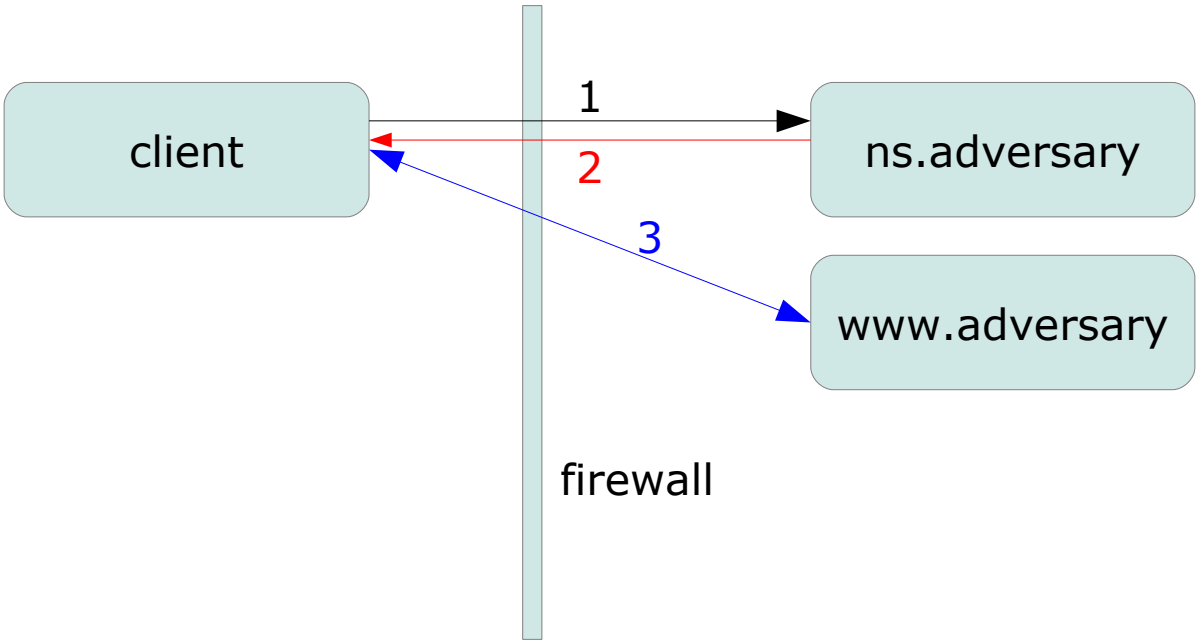

DNS Rebinding (cont'd)

- (1) is query for www.adversary
- (2) is a correct response pointing to (3), but with a short TTL
- Client queries again, but now get an IP to (4)

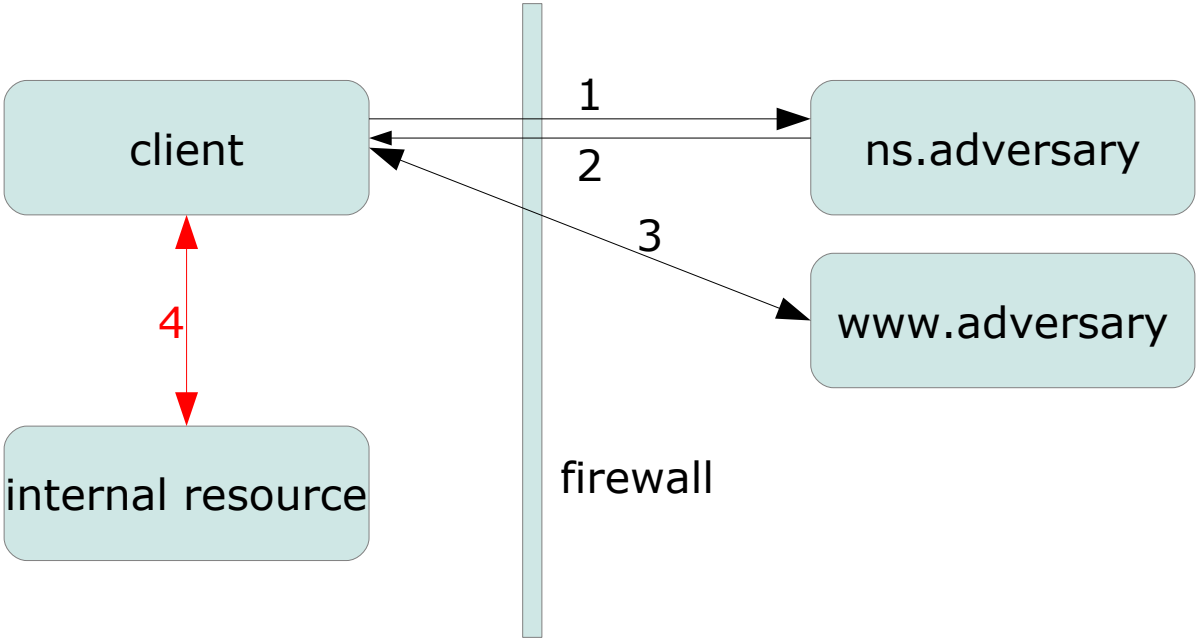

DNS Rebinding (cont'd)

- (1) is query for www.adversary
- (2) is a correct response pointing to (3), but with a short TTL
- Client queries again, but now get an IP to (4)

- Circumvents the same-origin policy
  - See the lecture on web security

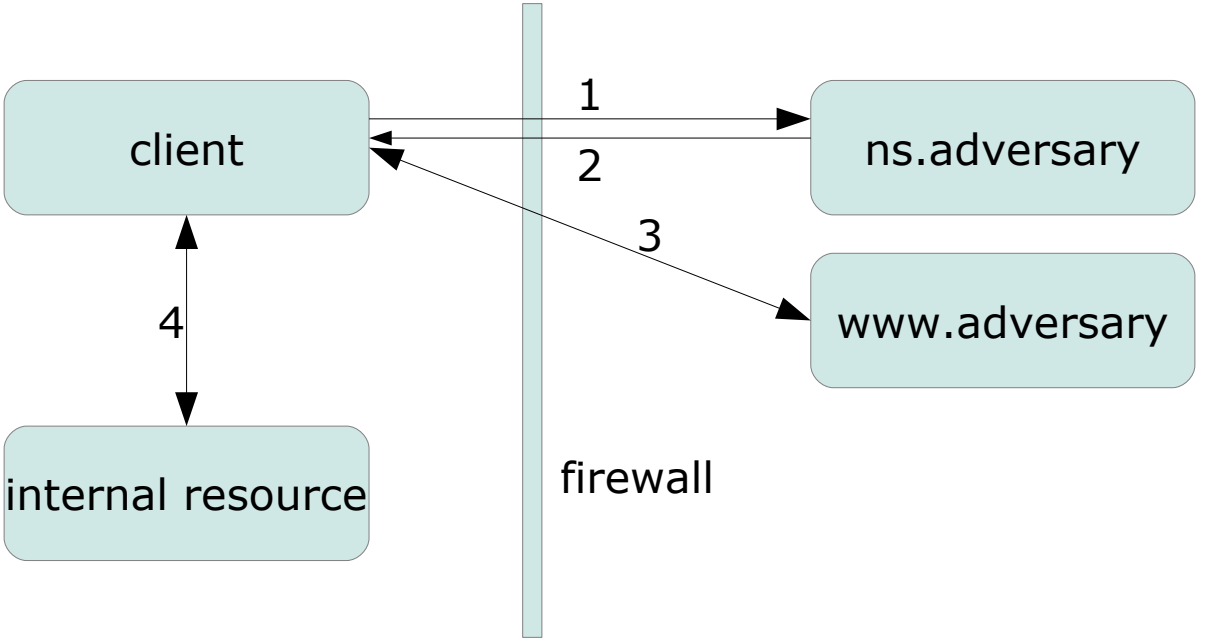

ROYAL INSTITUTE OF TECHNOLOGY

DNS Rebinding (cont'd)

- Browser mitigation: DNS Pinning
  - Refuse to switch to a new IP
  - But various services does this in a legitimate way
    - Proxies, VPN, dynamic DNS, ...
- Server-side defenses
  - Authenticate users with something other than IP
  - Reject HTTP requests with an unrecognized host header
- Firewall defenses
  - Check for external names resolving to internal addresses

ROYAL INSTITUTE OF TECHNOLOGY

Domain Name System Security Extensions (DNSSEC)

- Same principle as DNS
- All answers are digitally signed, to provide authentication
- New resource records (see RFC4034)
  - RRSIG DNS
    - Digital signature on the resource records in the response
  - DNSKEY
    - Public key that corresponds to the private used for the RRSIG
  - DS: Delegation Signer
    - Authenticate the DNSKEY record, i.e., a sub-domain
  - NSEC, to prove that some sub domains do not exist

ROYAL INSTITUTE OF TECHNOLOGY

DNSSEC (cont'd)

- Responses are not encrypted, i.e., no confidentiality
- Challenges as those for deploying a PKI
  - Centralized trust
  - Certificate revocation

ROYAL INSTITUTE OF TECHNOLOGY

DNSSEC NSEC

- Denial of existence records contains Next Secure (NSEC) resource records
- "NSEC RRs **assert which names do not exist** in a zone by linking from existing name to existing name along a canonical ordering of all the names within a zone." RFC4033

ROYAL INSTITUTE OF TECHNOLOGY

DNSSEC NSEC (cont'd)

- Allows for zone walking, i.e., zone enumeration
  - 1) Query for <random>.domain.com
  - 2) If domain exits, store it and repeat 1
  - 3) Get two valid sub-domain names, say a and b, store them
  - 4) Repeat step 1 for b1.domain.com
  - 5) When the complete linked list is found, the entire domain is mapped
- Potentially exposing servers not meant for public use
  - E.g., counteract the trouble of finding IPv6 hosts in the vast number space

ROYAL INSTITUTE OF TECHNOLOGY

DNSSEC NSEC3

- Hashed Authenticated Denial of Existence (NSEC3)
- Owner names are
  - Hashed, in order to hide them
  - Chained in hash order
- Still possible to zone walk to enumerate hosts
- See a discussion of NSEC/NSEC3 (which are mutually exclusive) here:
  http://www.internetsociety.org/deploy360/resources/dnssec-nsec-vs-nsec3/

ROYAL INSTITUTE OF TECHNOLOGY

DNSSEC Deployment

- 2010-07-15: Distribution of possible-to-validate signed root zone; publication of root zone trust anchor
- Domain owners, ISPs, and end users have been slow to adopt
- Deployment maps:
  - http://www.internetsociety.org/deploy360/dnssec/maps/

ROYAL INSTITUTE OF TECHNOLOGY

Summary

- DNS is a tree, with delegated trust
- DNS was not designed with security in mind
- Query IDs does not provide adequate protection
- DNSSEC addresses issues, but it is not everywhere just yet

ROYAL INSTITUTE OF TECHNOLOGY

Extra reading

- Request for Comments
  - RFC3833: Threat Analysis of the Domain Name System (DNS)
  - RFC4033: DNS Security Introduction and Requirements
  - RFC4034: Resource Records for the DNS Security Extension
- *An Illustrated Guide to the Kaminsky DNS Vulnerability*
  - http://unixwiz.net/techtips/iguide-kaminsky-dns-vuln.html